

# SDSC Summer Institute 2020

*Title: Parallel Computing Using MPI and OpenMP*

*Instructors: Mahidhar Tatineni*

*Date: August 6, 2020*

*Location: Main Room Session*

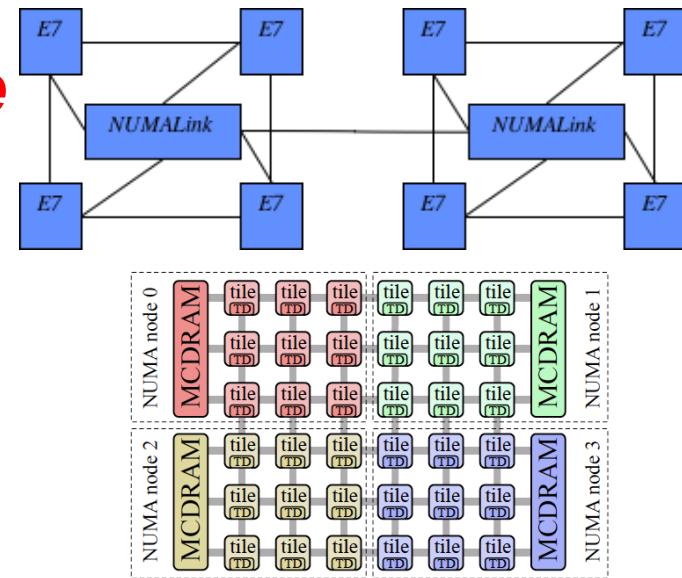
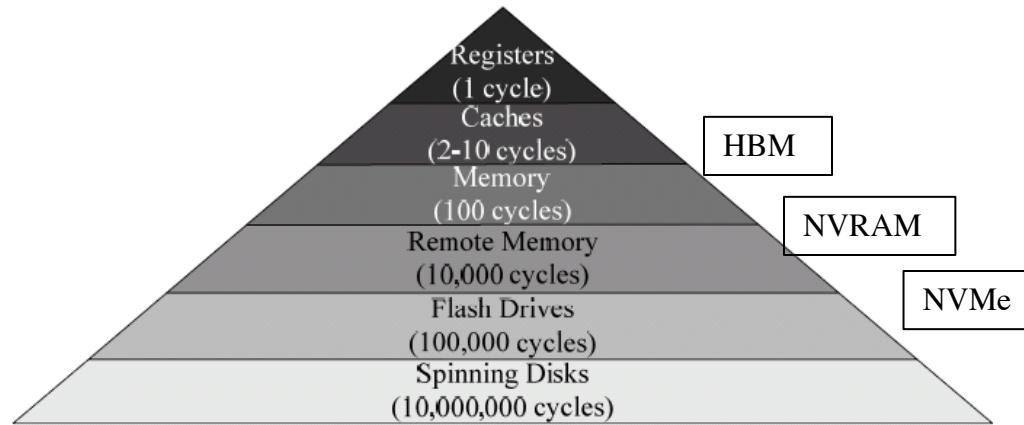
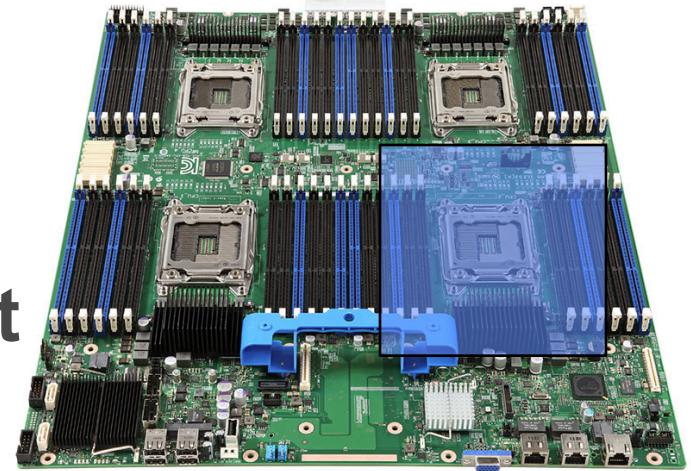
*Gitter (session support): Main Session*

*Gitter (general system support): Help Desk*



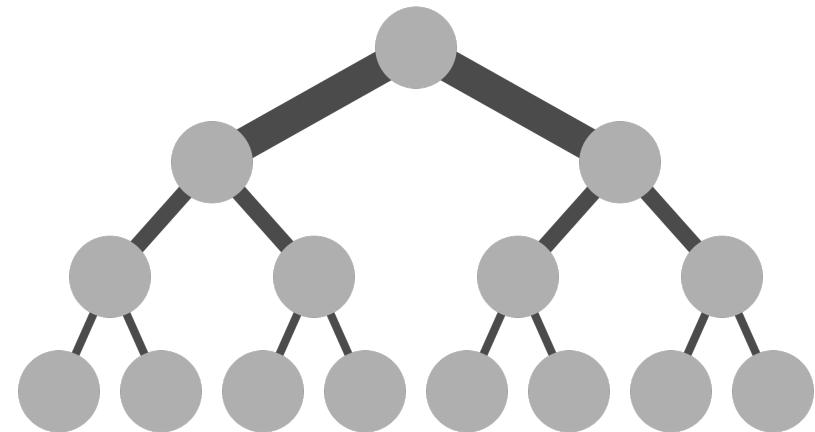
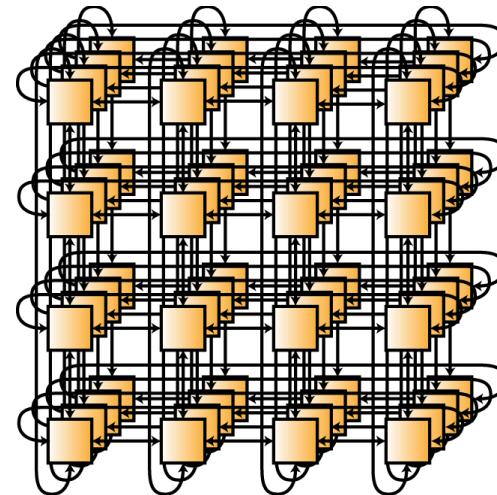
# Current Supercomputer Architectures

- Multi-socket server nodes
  - NUMA
  - Accelerators
- High performance interconnect
  - e.g. Infiniband, OmniPath
- ***Scalable parallel approach needed to achieve performance***



# Network Topologies

- Mesh, Torus, Hypercube
- Tree based
  - Fat-tree
  - Clos
- Dragonfly
- Metrics
  - Bandwidth
  - Diameter, Connectivity
  - Bisection bandwidth



# Parallel Computing

- Executing instructions concurrently on physical resources (not time slicing)
  - Multiple tightly coupled resources (e.g. cores) collaboratively solving a single problem
- Benefits
  - Capacity
    - Memory, storage
  - Performance
    - More instructions per unit of time (FLOPS)
    - Data streaming capability
- Cost and Complexity
  - Coordinate tasks and resources
  - Use resources efficiently

# Classification - Flynn's Taxonomy

Single Instruction Single Data	Single Instruction Multiple Data
Multiple Instruction Single Data	Multiple Instruction Multiple Data

- Single Program Multiple Data (SPMD)
- Multiple Program Multiple Data (MPMD)

- **Single Instruction, Single Data [Serial codes]**
- **Single Instruction, Multiple Data**
  - Processors run the same instructions, each operates on different data
  - Technically, Hadoop MapReduce fits this mode
  - GPUs
- **Multiple Instruction, Single Data**
  - Multiple instructions acting on single data stream. E.g. Different analysis on same set of data.
- **Multiple Instruction, Multiple Data**
  - Every processor may execute different instructions
  - Every processor may work on different parts of data
  - Execution can be synchronous or asynchronous, deterministic or non-deterministic

# **Memory, Communication, and Execution Models**

- **Shared**
  - Communication model: shared memory
- **Distributed**
  - Communication model: exchange messages
- **Execution Models**
  - Fork-Join (e.g. Thread Level Parallelism)
  - Single Program Multiple Data (SPMD)
- **Parallelism enabled by decomposing work**
  - Tasks can be executed concurrently
  - Some tasks can have dependencies

# What is OpenMP?

- **High level parallelism abstraction based on thread**
  - Easy to use
  - Suitable to an incremental approach
- **A specification and evolving standard**
  - “a portable, scalable model ... for developing portable parallel programs”
  - <http://openmp.org>
  - GNU, Intel, PGI, etc.
- **A set of**
  - Compiler directives
  - Library routines
  - Environment variables
  - Supports C/C++ and Fortran

```
#pragma omp parallel  
{  
....  
}
```

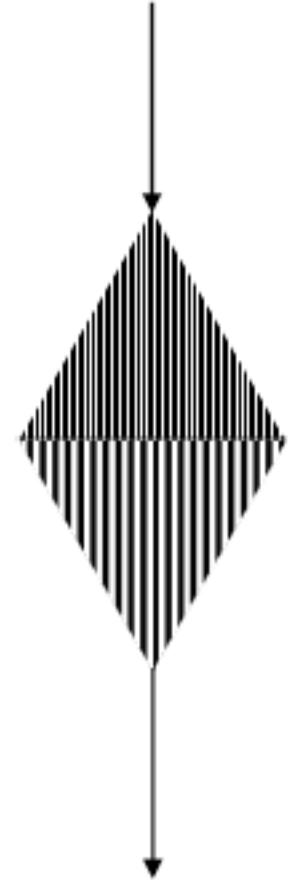
# OpenMP Models

- **Fork/Join Execution**

- Process starts single threaded (master thread)
- Forks child threads activated in parallel regions (team)
- The team synchronizes and threads are disbanded
  - – barrier
- Overhead is mitigated by reusing threads
- Master thread continues execution of serial phases

- **Work decomposition**

- Programming constructs
- Scope and compound statements
- Declarative in loops
- Mapping to threads can be static or dynamic
- Barriers and synchronization automatically inserted



# Compiler Directives

- Compiler directives is the main mechanism for introducing parallelism. Functionality enabled includes:
  - Spawning a parallel region
  - Dividing code among threads
  - Distributing loop iterations over threads
  - Serialization of parts of the code
  - Synchronization of work
- Example:  
**#pragma omp parallel default(shared) private(beta,pi)**

# Regions, Loops, Sections etc

```
#pragma omp parallel [clause[ [, ]clause] ...] new-line
structured-block
clause:
if(scalar-expression)
num_threads(integer-expression)
default(shared | none)
private(list)
firstprivate(list)
shared(list)
copyin(list)
reduction(operator: list)
```

```
#pragma omp for [clause[ [, ]clause] ...] new-line for-loops
clause:
private(list)
firstprivate(list)
lastprivate(list)
reduction(operator: list)
schedule(kind[, chunk_size])
collapse(n)
ordered
nowait
```

- #pragma omp single/master
- simd
- tasks

```
#pragma omp sections [clause[ [, ]clause] ...] new-line
{
#pragma omp section
structured-block
...
}
clause:
private(list)
firstprivate(list)
lastprivate(list)
reduction(operator: list)
nowait
```

# Scope of Variables

- **Clauses determine the scope of variables**
  - Default: shared (external)
  - Private
    - Also if declared inside region
  - firstprivate
  - shared
  - lastprivate
  - reductions
  - default
  - ....
- **Avoid race conditions!**

# Parallel Region Construct

```
!$OMP PARALLEL [clause ...]
    IF (scalar_logical_expression)
    PRIVATE (list)
    SHARED (list)
    DEFAULT (PRIVATE | FIRSTPRIVATE | SHARED | NONE)
    FIRSTPRIVATE (list)
    REDUCTION (operator: list)
    COPYIN (list)
    NUM_THREADS (scalar-integer-expression)
```

code block

```
!$OMP END PARALLEL
```

# Number of Threads

- Number of threads will be determined in the following order of precedence:
  - Evaluation of the IF clause
  - Setting of NUM\_THREADS clause
  - `omp_set_num_threads()` library function
  - `OMP_NUM_THREADS` environment variable
  - Default – usually ends up being the **\*number of cores on the node\*** (!)
- The last factor can accidentally lead to oversubscription of nodes in hybrid MPI/OpenMP codes.

# Work – Sharing Constructs

Directive format (C version):

```
#pragma omp for [clause ...] newline
    schedule (type [,chunk])
    ordered
    private (list)
    firstprivate (list)
    lastprivate (list)
    shared (list)
    reduction (operator: list)
    collapse (n)
    nowait
```

**for\_loop**

# Work-Sharing

- **Schedule:**
  - Static – Loop iterations are statically divided (chunk or as close to even as possible)
  - Dynamic – Loop iterations are divided in size chunk, and dynamically scheduled among threads. When a thread finishes one chunk it is dynamically assigned another
  - Guided – Similar to dynamic but chunk size is proportionally reduced based on work remaining.
  - Runtime – set at runtime by environment variables
  - Auto – set by compiler or runtime system.

# Simple OpenMP Program – Compute PI

- Find the number of tasks and taskids (`omp_get_num_threads`, `omp_get_thread_num`)
- PI is calculated using an integral. The number of intervals used for the integration is fixed at 128000.
- Use OpenMP loop parallelization to divide up the compute work.
- Introduce concept of private and shared variables.
- OpenMP reduction operation used to compute the sum for the final integral.
- Today's OpenMP examples are in: **/home/\$USER/PARALLEL/OPENMP**  
(Run “make” in that directory to generate executables).

# OpenMP Program to Compute PI

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[])
{
    int nthreads, tid;
    int i, INTERVALS;
    double n_1, x, pi = 0.0;

    INTERVALS=128000;
    /* Fork a team of threads giving them their own
       copies of variables */
    #pragma omp parallel private(nthreads, tid)
    {
        /* Obtain thread number */
        tid = omp_get_thread_num();
        printf("Hello from thread = %d\n", tid);
        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
        /* All threads join master thread and disband */
        n_1 = 1.0 / (double)INTERVALS;

        /* Parallel loop with reduction for calculating PI */
        #pragma omp parallel for private(i,x)
        shared(n_1,INTERVALS) reduction(+:pi)
        for (i = 0; i < INTERVALS; i++)
        {
            x = n_1 * ((double)i - 0.5);
            pi += 4.0 / (1.0 + x * x);
        }
        pi *= n_1;
        printf ("Pi = %.12lf\n", pi);
    }
}
```

# OpenMP result: 1-D Heat Equation

```
$ sbatch --res=si2020resday4 pi_openmp.sb
```

```
[train129@comet-ln2 OPENMP]$ more pi.18444806.comet-06-08.out
Hello from thread = 8
Hello from thread = 0
Hello from thread = 3
Hello from thread = 2
Hello from thread = 15
Hello from thread = 1
Hello from thread = 11
Hello from thread = 5
Hello from thread = 12
Hello from thread = 4
Hello from thread = 9
Hello from thread = 6
Hello from thread = 7
Hello from thread = 10
Hello from thread = 13
Hello from thread = 14
Threre are 16 threads!
PI = 3.139384000000, Err = 2.208654e-03, Time = 0.045058
[train129@comet-ln2 OPENMP]$
```

# More Work-Share Constructs

- **SECTIONS** directive – enclosed sections are divided among the threads.
- **WORKSHARE** directive – divides execution of block into units of work, each of which is executed once.
- **SINGLE** directive – Enclosed code is executed by only one thread.

# Synchronization Constructs

- **MASTER** directive – Specifies region is executed only by the master thread.
- **CRITICAL** directive – Region of the code that is executed one thread at a time.
- **BARRIER** directive – synchronize all threads
- **TASKWAIT** directive – wait for all child tasks to complete
- **ATOMIC** directive – specific memory location updated atomically (not let all threads write at the same time)

# Simple Application using OpenMP: 1-D Heat Equation

- $\partial T / \partial t = \alpha (\partial^2 T / \partial x^2); T(0) = 0; T(1) = 0; (0 \leq x \leq 1)$   
 $T(x, 0)$  is known as an initial condition.
- Discretizing for numerical solution we get:  
$$T^{(n+1)}_i - T^{(n)}_i = (\alpha \Delta t / \Delta x^2)(T^{(n)}_{i-1} - 2T^{(n)}_i + T^{(n)}_{i+1})$$
  
( $n$  is the index in time and  $i$  is the index in space)

# Fortran OpenMP Code: 1-D Heat Equation

```
PROGRAM HEATEQN
```

```
implicit none
integer :: iglobal, itime, nthreads
real*8 :: xalp, delx, delt, pi
real*8 :: T(0:100,0:10)
integer::: id
integer::: OMP_GET_THREAD_NUM,
          OMP_GET_NUM_THREADS

!$OMP PARALLEL SHARED(nthreads)
!$OMP MASTER
nthreads = omp_get_num_threads()
write (*,*) 'There are', nthreads, 'threads'
!$OMP END MASTER
!$OMP END PARALLEL
if (nthreads.ne.3) then
  write(*,*)"Use exactly 3 threads for this case"
  stop
endif
delx = 0.1d0
delt = 1d-4
xalp = 2.0d0
```

```
***** Initial Conditions *****
pi = 4d0*datan(1d0)
do iglobal = 0, 10
  T(0,iglobal) = dsin(pi*delx*dfloat(iglobal))
enddo
***** Iterations *****
do itime = 1 , 3
  write(*,*)"Running Iteration Number ", itime
!$OMP PARALLEL DO PRIVATE(iglobal)
SHARED(T,xalp,delx,delt,itime)
  do iglobal = 1, 9
    T(itime,iglobal)=T(itime-1,iglobal)+  

    + xalp*delt/delx*  

    + (T(itime-1,iglobal-1)-2*T(itime-1,iglobal)+T(itime-  

    1,iglobal+1))
  enddo
!$OMP BARRIER
  enddo
  do iglobal = 0, 10
    write(*,*)iglobal,T(3, iglobal)
  enddo
END
```

# OpenMP result: 1-D Heat Equation

```
$ sbatch --res=si2020resday4 heat_openmp.sb
```

*Sample output file:*

```
There are      3 threads
Running Iteration Number      1
Running Iteration Number      2
Running Iteration Number      3
0 0.000000000000000E+000
1 0.307205621017285
2 0.584339815421976
3 0.804274757358271
4 0.945481682332598
5 0.994138272681972
6 0.945481682332598
7 0.804274757358271
8 0.584339815421977
9 0.307205621017285
10 7.797843424221369E-316
```

# Data Scope Attribute Clauses

- **PRIVATE** – variables in the list are private to each thread.
- **SHARED** – variables in the list are shared between all threads.
- **DEFAULT** – default scope for all variables in a parallel region.
- **FIRSTPRIVATE** – variables are private and initialized according to value prior to entry into parallel or work sharing construct.
- **LASTPRIVATE** – variables are private, the value from the last iteration or section is copied to original variable object.
- **Others** – **COPYIN**, **COPYPRIVATE**
- **REDUCTION** – reduction on variables in the list

# **Run Time Library Routines**

- **Setting and querying number of threads**
- **Querying thread identifier, team size**
- **Setting and querying dynamic threads feature**
- **Querying if in parallel region and at what level**
- **Setting and querying nested parallelism**
- **Setting, initializing and terminating locks, nested locks.**
- **Querying wall clock time and resolution.**

# Environment Variables

- **OMP\_SCHEDULE** – e.g set to “dynamic”
- **OMP\_NUM\_THREADS**
- **OMP\_DYNAMIC** (TRUE or FALSE)
- **OMP\_PROC\_BIND** (TRUE or FALSE)
- **OMP\_NESTED** (TRUE or FALSE)
- **OMP\_STACKSIZE** – size of stack for created threads
- **OMP\_THREAD\_LIMIT**

# General OpenMP Performance Considerations

- Avoid or minimize use of BARRIER, CRITICAL (complete serialization here!), ORDERED regions, and locks. Can use NOWAIT clause to avoid redundant barriers.
- Parallelize at a high level, i.e. maximize the work in the parallel regions to reduce parallelization overhead.
- Use appropriate loop scheduling – static has low synchronization overhead but can be unbalanced, dynamic (and guided) have higher synchronization overheads but can improve load balancing.
- Avoid false sharing (more about it in following slide)!

# What is False Sharing?

- Most modern processors have a cache buffer between slow memory and high speed registers of the CPU.
- Accessing a memory location causes a “cache line” to be copied into the cache.
- In an OpenMP code two processors may be accessing two different elements in the same cache line. On writes this will lead to “cache line” being marked invalid (because cache coherency is being maintained).
- This will lead to an increase in memory traffic even though the write is to different elements (hence the term **false sharing**).
- This can have a drastic performance impact if such updates are occurring frequently in a loop.

# False Sharing Example

Code snippet:

```
double global=0.0, local[NUM_THREADS];
#pragma omp parallel num_threads(NUM_THREADS)
{
    int tid = omp_get_thread_num();
    local[tid] = 0.0;
    #pragma omp for
    for (i = 0; i < N; i++)
        local[tid] += x[i];
    #pragma omp atomic
    global += local[me];
}
```

# False Sharing - Solutions

- Three options
  - Compiler directives to align individual variables on cache line boundaries

```
_declspec (align(64)) int thread1_global_variable;  
_declspec (align(64)) int thread2_global_variable;
```
  - Pad arrays/data structures to make sure array elements begin on cache line boundary.
  - Use thread local copies of data (assuming the copy overhead is small compared to overall run time).

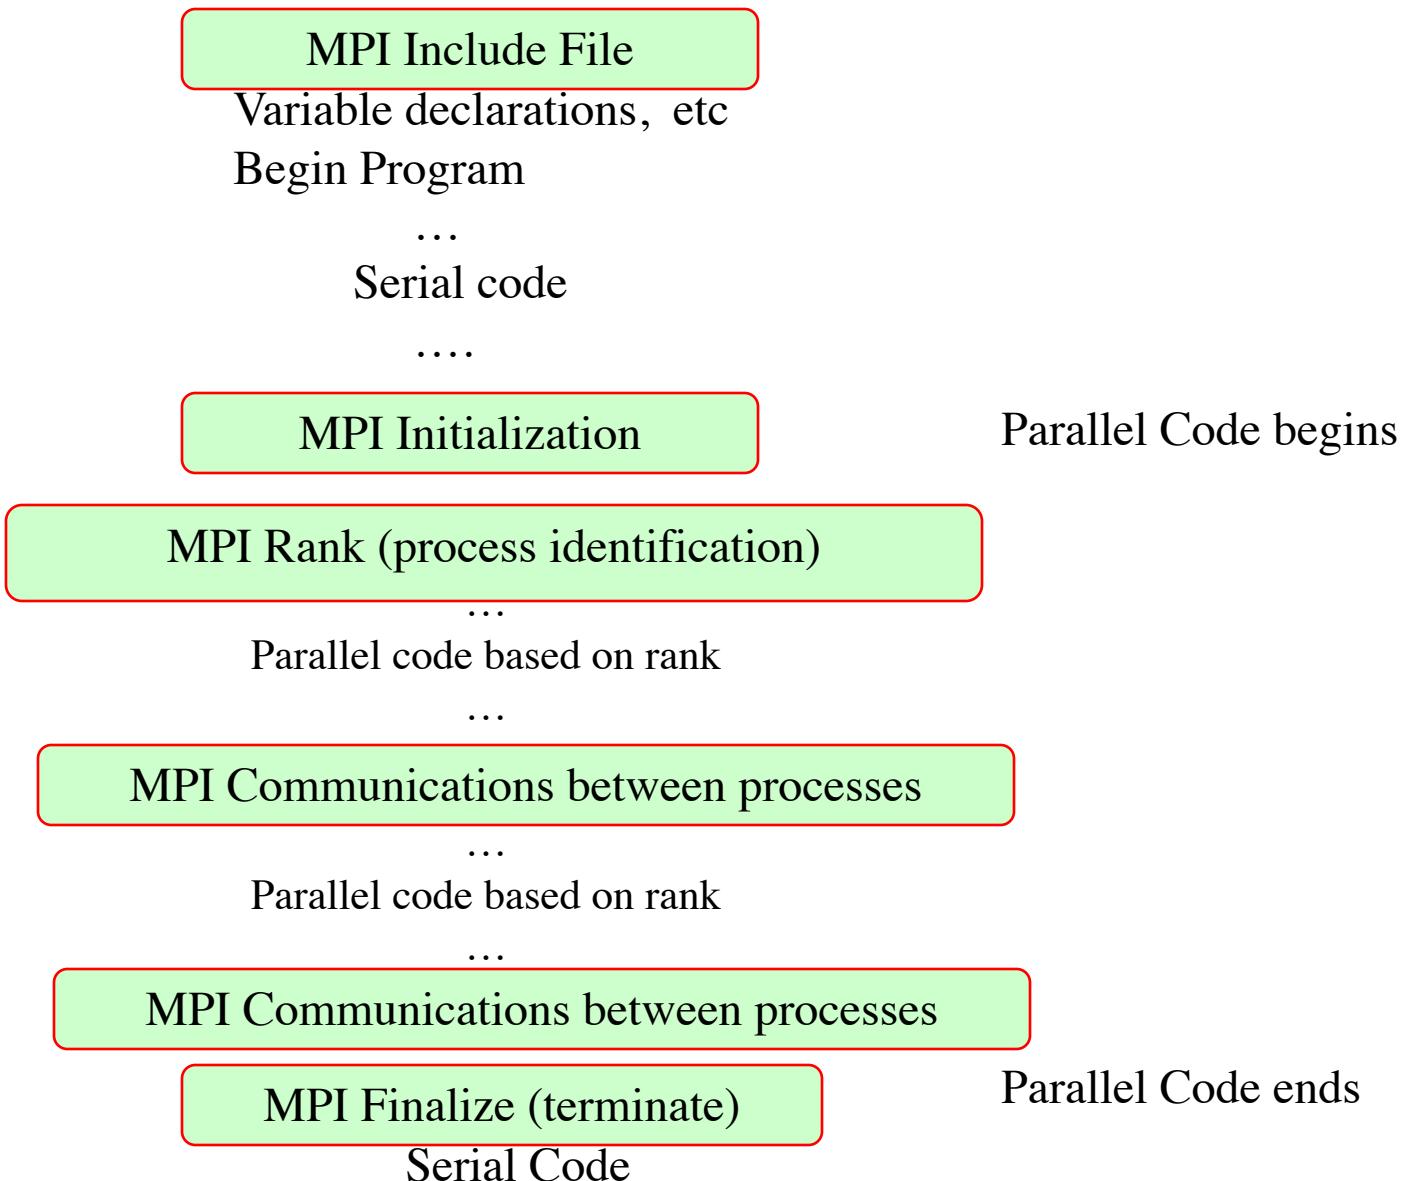
# Homework!

- Download matrix multiply example from LLNL site:
  - [https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp\\_mm.f](https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp_mm.f) (wget  
[https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp\\_mm.f](https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp_mm.f) on Comet)
- Compile (ifort -fopenmp omp\_mm.f) and run the example. See if you can vary the environmental variables, scheduling to get better performance!
- This is very quick intro. Lot of ongoing developments. Detailed specifications at:
  - <https://www.openmp.org/specifications/>

# Message Passing Interface (MPI)

- **Low level message passing abstraction**
  - SPMD execution model + messages
  - Designed for distributed memory. Implemented on hybrid distributed memory/shared memory systems.
- **MPI: API specification**
  - Portable: de-fact standard for parallel computing, portable, system specific optimizations without changing code interface
  - <http://www mpi-forum.org>
  - Several implementations - e.g MVAPICH2 (default on Comet), Intel MPI, and OpenMPI
  - High performance implementations available virtually on any interconnect and system
  - Point-to-point communication, datatypes, collective operations
  - One-sided communication, Parallel file I/O, Tool support, ...

# Typical MPI Code Structure



# Simple MPI Program – Compute PI

- Initialize MPI (`MPI_Init` function)
- Find the number of tasks and taskids (`MPI_Comm_size`, `MPI_Comm_rank`)
- PI is calculated using an integral. The number of intervals used for the integration is fixed at 128000.
- Computes the sums for a different sections of the intervals in each MPI task.
- At the end of the code, the sums from all the tasks are added together to evaluate the final integral. This is accomplished through a reduction operation (`MPI_Reduce` function).
- Simple code illustrates decomposition of problem into parallel components.

# MPI Program to Compute PI

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char *argv[])
{
    int numprocs, rank;
    int i, iglob, INTERVALS, INTLOC;
    double n_1, x;
    double pi, piloc;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD,
                  &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    INTERVALS=128000;
    printf("Hello from MPI task= %d\n", rank);
    MPI_Barrier(MPI_COMM_WORLD);
    if (rank == 0)
    {
        printf("Number of MPI tasks = %d\n", numprocs);
    }

    INTLOC=INTERVALS/numprocs;
    piloc=0.0;
    n_1=1.0/(double)INTERVALS;
    for (i = 0; i < INTLOC; i++)
    {
        iglob = INTLOC*rank+i;
        x = n_1 * ((double)iglob - 0.5);
        piloc += 4.0 / (1.0 + x * x);
    }

    MPI_Reduce(&piloc,&pi,1,MPI_DOUBLE,MPI_SUM,0,
               MPI_COMM_WORLD);
    if (rank == 0)
    {
        pi *= n_1;
        printf ("Pi = %.12lf\n", pi);
    }

    MPI_Finalize();
}
```

# PI Code : MPI Environment Functions

**MPI\_Init(&argc, &argv);**

Initializes MPI, \*must\* be called (only once) in every MPI program before any MPI functions.

**MPI\_Comm\_size(MPI\_COMM\_WORLD, &numprocs);**

Returns the total number of tasks in the communicator. MPI uses communicators to define which collections of processes can communicate with each other. The default MPI\_COMM\_WORLD includes all the processes. User defined communicators are an option.

**MPI\_Comm\_rank(MPI\_COMM\_WORLD, &rank);**

Returns the rank (ID) of the calling MPI process within the communicator.

**MPI\_Finalize();**

Ends the MPI execution environment. No MPI calls after this.!

The other routines in the code are collectives and we will discuss them later in the talk.

# Compiling and Running PI Example

`cd /home/$USER/PARALLEL/SIMPLE`

Compile: `mpicc -o pi_mpi.exe pi_mpi.c`

Submit Job: `sbatch --res=si2020resday4 pi_mpi.sb`

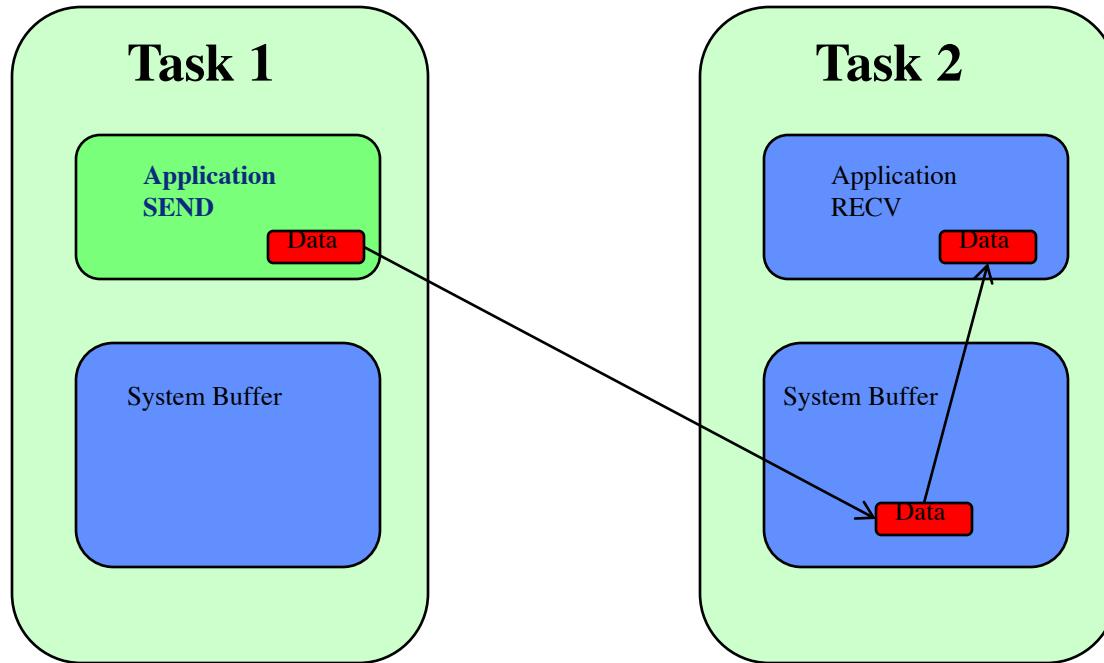
*Sample Output:*

```
Hello from MPI task= 12
Hello from MPI task= 14
Hello from MPI task= 8
Hello from MPI task= 3
Hello from MPI task= 2
Hello from MPI task= 4
Hello from MPI task= 13
Hello from MPI task= 9
Hello from MPI task= 5
Hello from MPI task= 1
Hello from MPI task= 11
Hello from MPI task= 10
Hello from MPI task= 15
Hello from MPI task= 7
Hello from MPI task= 6
Hello from MPI task= 0
Number of MPI tasks = 16
Pi = 3.141594606714
```

# Point to Point Communication

- Passing data between two, and only two different MPI tasks.
- Typically one task performs a send operation and the other task performs a matching receive.
- MPI Send operations have choices with different synchronization (when does a send complete) and different buffering (where the data resides till it is received) modes.
- Any type of send routine can be paired with any type of receive routine.
- MPI also provides routines to probe status of messages, and “wait” routines.

# Buffers



- Buffer space is used for data in transit – whether its waiting for a receive to be ready or if there are multiple sends arriving at the same receiving tasks.
- Typically a system buffer area managed by the MPI library (opaque to the user) is used. Can exist on both sending & receiving side.
- MPI also provides for user managed send buffer.

# Blocking MPI Send, Receive Routines

- Blocking send call will return once it is safe for the application buffer (send data) to be reused.
- This can happen as soon as the data is copied into the system (MPI) buffer on receiving process.
- Synchronous if there is confirmation of safe send, and asynchronous otherwise.
- Blocking receive returns once the data is in the application buffer (receive data) and can be used by the application.

# Blocking Send, Recv Example (Code Snippet)

```
if(myid == 0) {  
    for(i = 0; i < 10; i++) {  
        s_buf[i] = i*4.0;  
    }  
    MPI_Send(s_buf, size, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);  
}  
else if(myid == 1) {  
    MPI_Recv(r_buf, size, MPI_FLOAT, 0, tag, MPI_COMM_WORLD,  
&reqstat);  
    for (i = 0; i < 10; i++ ){  
        printf("r_buf[%d] = %f\n", i, r_buf[i] );  
    }  
}
```

# Blocking Send, Recv Example

Location: \$HOME/PARALLEL/PTOP

Compile: **mpicc -o blocking.exe blocking.c**

Submit Job: **sbatch --res=si2020resday4 blocking.sb**

Output:

r\_buf[0] = 0.000000

r\_buf[1] = 4.000000

r\_buf[2] = 8.000000

r\_buf[3] = 12.000000

r\_buf[4] = 16.000000

r\_buf[5] = 20.000000

r\_buf[6] = 24.000000

r\_buf[7] = 28.000000

r\_buf[8] = 32.000000

r\_buf[9] = 36.000000

# Deadlocking MPI Tasks

- Take care to sequence blocking send/recvs. Easy to deadlock processes waiting on each other with circular dependencies.
- Can also occur with control errors and unexpected semantics
- For example, take the following code snippet:

```
if(myid == 0) {  
    MPI_Ssend(s_buf, size, MPI_FLOAT, 1, tag1, MPI_COMM_WORLD);  
    MPI_Recv(r_buf, size, MPI_FLOAT, 1, tag2, MPI_COMM_WORLD, &reqstat);  
}  
else if(myid == 1) {  
    MPI_Ssend(s_buf, size, MPI_FLOAT, 0, tag2, MPI_COMM_WORLD);  
    MPI_Recv(r_buf, size, MPI_FLOAT, 0, tag1, MPI_COMM_WORLD, &reqstat);  
    for (i = 0; i < 10; i++) {  
        printf("r_buf[%d] = %f\n", i, r_buf[i]);  
    }  
}
```

- The MPI\_Ssend on both tasks will not complete till the MPI\_Recv is posted (which will never happen given the order).

# Deadlock Example

- Location: \$HOME/PARALLEL/PTOP
- Compile: `mpicc -o deadlock.exe deadlock.c`
- Submit Job: `sbatch --res=si2020resday4 deadlock.sb`
- It should technically finish in less than a second since the data transferred is a few bytes. However, the code deadlocks and hits the wallclock limit (2 minutes in the script).
- Error info:

`slurmstepd: *** JOB 35124016 ON comet-03-10 CANCELLED AT 2020-08-05T19:03:43 DUE TO TIME LIMIT ***`

`[comet-03-10.sdsc.edu:mpirun_rsh][signal_processor] Caught signal 15, killing job  
Connection to comet-03-10 closed by remote host.`

# Deadlock Example – Simple Fix

- Change the order on one of processes!
- For example, take the following code snippet:

```
if(myid == 0) {  
    MPI_Ssend(s_buf, size, MPI_FLOAT, 1, tag1, MPI_COMM_WORLD);  
    MPI_Recv(r_buf, size, MPI_FLOAT, 1, tag2, MPI_COMM_WORLD, &reqstat);  
}  
else if(myid == 1) {  
    MPI_Recv(r_buf, size, MPI_FLOAT, 0, tag1, MPI_COMM_WORLD, &reqstat);  
    MPI_Ssend(s_buf, size, MPI_FLOAT, 0, tag2, MPI_COMM_WORLD);  
    for (i = 0; i < 10; i++) {  
        printf("r_buf[%d] = %f\n", i, r_buf[i]);  
    }  
}
```

- Now the MPI\_Ssend on task 0 will complete since the corresponding MPI\_Recv is posted first on task 1. (qsub deadlock-fix1.cmd)
- We will look at **Non-Blocking** options next.

# Deadlock Example (Fix 1)

- Location: \$HOME/PARALLEL/PTOP
- Compile: **mpicc -o deadlock-fix1.exe deadlock-fix1.c**
- Submit Job: **sbatch --res=si2020resday4 deadlock-fix1.sb**
- **Fix works!**

```
$ more deadlock-fix1.out
r_buf[0] = 0.000000
r_buf[1] = 4.000000
r_buf[2] = 8.000000
r_buf[3] = 12.000000
r_buf[4] = 16.000000
r_buf[5] = 20.000000
r_buf[6] = 24.000000
r_buf[7] = 28.000000
r_buf[8] = 32.000000
r_buf[9] = 36.000000
```

# Non-Blocking MPI Send, Receive Routines

- Non-Blocking MPI Send, Receive routines return before there is any confirmation of receives or completion of the actual message copying operation.
- The routines simply put in the request to perform the operation.
- MPI wait routines can be used to check status and block till the operation is complete and it is safe to modify/use the information in the application buffer.
- This non-blocking approaches allows computations (that don't depend on this data in transit) to continue while the communication operations are in progress. This allows for hiding the communication time with useful work and hence improves parallel efficiency.

# Non-Blocking Send, Recv Example

- Example uses **MPI\_Isend**, **MPI\_Irecv**, **MPI\_Wait**
- **Code snippet:**

```
if(myid == source){  
    s_buf=1024;  
    MPI_Isend(&s_buf,count,MPI_INT,destination,tag,MPI_COMM_WORLD,&request);  
}  
if(myid == destination {  
    MPI_Irecv(&r_buf,count,MPI_INT,source,tag,MPI_COMM_WORLD,&request);  
}  
MPI_Wait(&request,&status);
```

- **Compile & Run:**

```
mpicc -o nonblocking.exe nonblocking.c  
sbatch --res=si2020resday4 nonblocking.sb
```

*Sample output:*

```
processor 0 sent 1024  
processor 1 got 1024
```

# Deadlock Example – Non-Blocking Option

- Change the order on one of processes!
- For example, take the following code snippet:

```
if(myid == 0) {  
    MPI_Isend(s_buf, size, MPI_FLOAT, 1, tag1, MPI_COMM_WORLD, &request);  
    MPI_Recv(r_buf, size, MPI_FLOAT, 1, tag2, MPI_COMM_WORLD, &reqstat);  
}  
else if(myid == 1) {  
    MPI_Ssend(s_buf, size, MPI_FLOAT, 0, tag2, MPI_COMM_WORLD);  
    MPI_Recv(r_buf, size, MPI_FLOAT, 0, tag1, MPI_COMM_WORLD, &reqstat);  
    for (i = 0; i < 10; i++) {  
        printf("r_buf[%d] = %f\n", i, r_buf[i]);  
    }  
}
```

- Now the MPI\_Ssend on task 0 will complete since the corresponding MPI\_Recv is posted first on task 1. (qsub deadlock-fix1.cmd)
- We will look at **Non-Blocking** options next.

# Deadlock Example (Fix 2)

- Location: \$HOME/PARALLEL/PTOP
- Compile: **mpicc -o deadlock-fix2-nb.exe deadlock-fix2-nb.c**
- Submit Job: **sbatch --res=si2020resday4 deadlock-fix2-nb.sb**
- **Fix works!**

```
$ more deadlock-fix2-nb.out
```

```
r_buf[0] = 0.000000
r_buf[1] = 4.000000
r_buf[2] = 8.000000
r_buf[3] = 12.000000
r_buf[4] = 16.000000
r_buf[5] = 20.000000
r_buf[6] = 24.000000
r_buf[7] = 28.000000
r_buf[8] = 32.000000
r_buf[9] = 36.000000
```

# Collective MPI Routines

- **Synchronization Routines:** All processes in group/communicator wait till they get synchronized.
- **Data Movement:** Send/Receive data from all processes.  
E.g. Broadcast, Scatter, Gather, AlltoAll.
- **Collective Computation (reductions):** Perform reduction operations (min, max, add, multiply, etc.) on data obtained from all processes.
- **Collective Computation and Data Movement combined (Hybrid).**

# Synchronization Example

- Our simple PI program had a synchronization example.
- Code Snippet:

```
printf("Hello from MPI task= %d\n", rank);
MPI_Barrier(MPI_COMM_WORLD);
if (rank == 0)
{
    printf("Number of MPI tasks = %d\n", numprocs);
}
```

- All tasks will wait till they are synchronized at this point.

# Broadcast Example

- **Code Snippet** (**All collectives examples in \$HOME/PARALLEL/COLLECTIVES**):
  - if(myid .eq. source)then
  - do i=1,count
  - buffer(i)=i
  - enddo
  - endif
  - **Call MPI\_Bcast(buffer, count, MPI\_INTEGER,source,& MPI\_COMM\_WORLD,ierr)**
- **Compile:**
  - **mpif90 -o bcast.exe bcast.f90**
- **Run:**
  - **sbatch --res=si2020resday4 bcast.sb**
- **Output:**

processor	1	got	1	2	3	4
processor	0	got	1	2	3	4
processor	2	got	1	2	3	4
processor	3	got	1	2	3	4

# Reduction Example

- **Code Snippet:**

```
myidp1 = myid+1
call MPI_Reduce(myidp1,ifactorial,1,MPI_INTEGER,MPI_PROD,root,MPI_COMM_WORLD,ierr)
if (myid.eq.root) then
    write(*,*)numprocs,"! = ",ifactorial
endif
```

- **Compile:**

```
mpif90 -o factorial.exe factorial.f90
```

- **Run:**

```
sbatch --res=si2020resday4 factorial.sb
```

- **Output:**

```
8 ! =      40320
```

# MPI\_Allreduce example

- Code Snippet:

```
imaxloc=IRAND(myid)
call MPI_ALLREDUCE(imaxloc,imax,1,MPI_INTEGER,MPI_MAX,MPI_COMM_WORLD,
mpi_err)
if (imax.eq.imaxloc) then
    write(*,*)"Max=",imax,"on task",myid
endif
• Compile:
  mpif90 -o allreduce.exe allreduce.f90
```

- Run:

```
  sbatch --res=si2020resday4 allreduce.sb
```

- Output:

```
  Max= 337897 on task      7
```

# Data Types

C Data Types	FORTRAN Data Types
MPI_CHAR	MPI_CHARACTER
MPI_WCHAR	MPI_INTEGER
MPI_SHORT	MPI_INTEGER1
MPI_INT	MPI_INTEGER2
MPI_LONG	MPI_INTEGER4
MPI_LONG_LONG_INT	MPI_REAL
MPI_LONG_LONG	MPI_REAL2
MPI_SIGNED_CHAR	MPI_REAL4
MPI_UNSIGNED_CHAR	MPI_REAL8
MPI_UNSIGNED_SHORT	MPI_DOUBLE_PRECISION
MPI_UNSIGNED_LONG	MPI_COMPLEX
MPI_UNSIGNED	MPI_DOUBLE_COMPLEX
MPI_FLOAT	MPI_LOGICAL
MPI_DOUBLE	MPI_BYTE
MPI_LONG_DOUBLE	MPI_PACKED
MPI_C_COMPLEX	
MPI_C_FLOAT_COMPLEX	

# MPI Reduction Operations

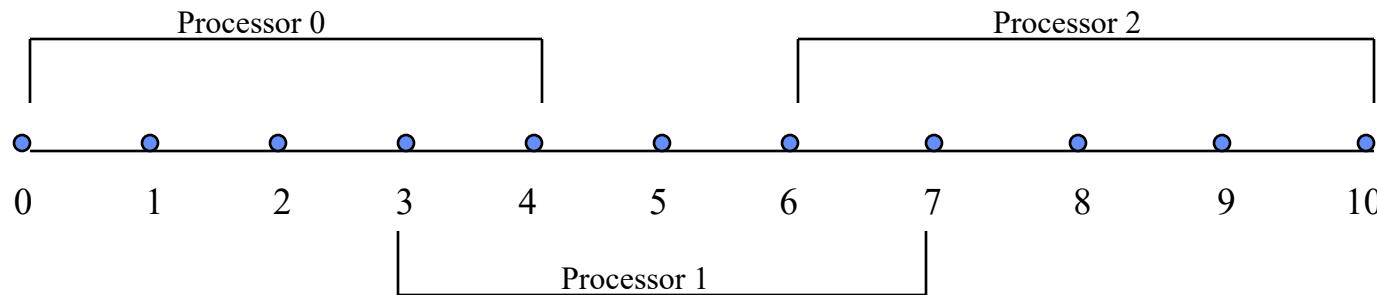
NAME	OPERATION
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical AND
MPI_BAND	Bit-wise AND
MPI_LOR	Logical OR
MPI_BOR	Bit-wise OR
MPI_LXOR	Logical XOR
MPI_BXOR	Bit-wise XOR
MPI_MAXLOC	Maximum value and location
MPI_MINLOC	Minimum value and location

# Decomposition and Mapping

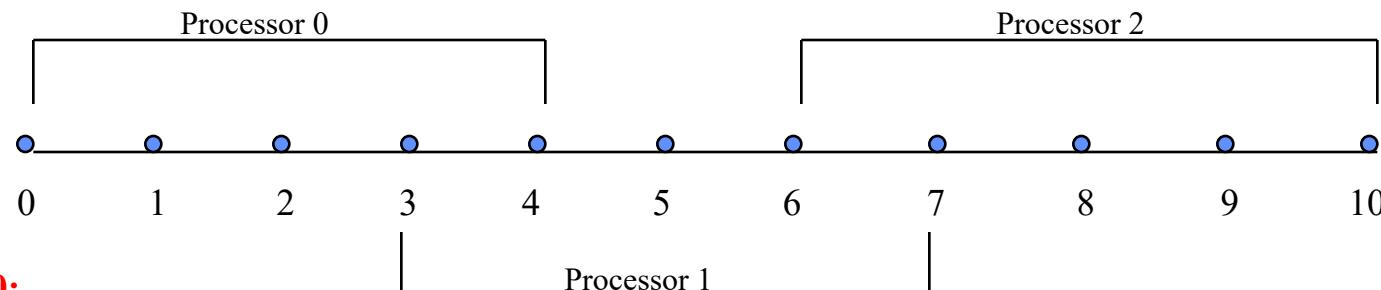
- **Data and work decomposition**
  - Map partitioned domain to processes
- **Mapping**
  - Processes/ranks topology
  - System/Domain/Data
- **How to share data?**
  - Exchange messages and replicate data
- **Load imbalance**
  - What if the system is not regular?
  - Is work proportional to size of partitions?

## Simple Application using MPI: 1-D Heat Equation

- $\partial T / \partial t = \alpha (\partial^2 T / \partial x^2)$ ;  $T(0) = 0$ ;  $T(1) = 0$ ; ( $0 \leq x \leq 1$ )  
 $T(x, 0)$  is known as an initial condition.
- Discretizing for numerical solution we get:  
$$T^{(n+1)}_i - T^{(n)}_i = (\alpha \Delta t / \Delta x^2)(T^{(n)}_{i-1} - 2T^{(n)}_i + T^{(n)}_{i+1})$$
  
( $n$  is the index in time and  $i$  is the index in space)
- In this example we solve the problem using 11 points and we distribute this problem over exactly 3 processors (for easy demo) shown graphically below:



# Simple Application using MPI: 1-D Heat Equation



## Processor 0:

Local Data Index : ilocal = 0 , 1, 2, 3, 4

Global Data Index: iglobal = 0, 1, 2, 3, 4

Solve the equation at (1,2,3)

**Data Exchange: Get 4 from processor 1; Send 3 to processor 1**

## Processor 1:

Local Data Index : ilocal = 0, 1, 2, 3, 4

Global Data Index : iglobal = 3, 4, 5, 6, 7

Solve the equation at (4,5,6)

**Data Exchange: Get 3 from processor 0; Get 7 from processor 2; Send 4 to processor 0; Send 6 to processor 2**

## Processor 2:

Local Data Index : ilocal = 0, 1, 2, 3, 4

Global Data Index : iglobal = 6, 7, 8, 9, 10

Solve the equation at (7,8,9)

**Data Exchange: Get 6 from processor 1; Send 7 to processor 1**

# FORTRAN MPI CODE: 1-D Heat Equation

## PROGRAM HEATEQN

```
implicit none
include "mpif.h"
integer :: iglobal, ilocal, itime
integer :: ierr, nnodes, my_id
integer :: dest, from, status(MPI_STATUS_SIZE),tag
integer :: msg_size
real*8 :: xalp,delx,delt,pi
real*8 :: T(0:100,0:5), TG(0:10)
CHARACTER(20) :: FILEN

delx = 0.1d0
delt = 1d-4
xalp = 2.0d0

call MPI_INIT(ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD,
nnodes, ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD,
my_id, ierr)

if (nnodes.ne.3) then
if (my_id.eq.0) then
print *, "This test needs exactly 3 tasks"
endif
```

```
print *, "Process ", my_id, "of", nnodes , "has started"
!***** Initial Conditions
*****  

pi = 4d0*datan(1d0)
do ilocal = 0, 4
iglobal = 3*my_id+ilocal
T(0,ilocal) = dsin(pi*delx*dfloat(iglobal))
enddo
write(*,*)"Processor", my_id, "has finished setting
+ initial conditions"
!***** Iterations
*****  

do itime = 1 , 3
if (my_id.eq.0) then
write(*,*)"Running Iteration Number ", itime
endif
do ilocal = 1, 3
T(itime,ilocal)=T(itime-1,ilocal)+  

+ xalp*delt/delx/delx*  

+ (T(itime-1,ilocal-1)-2*T(itime-1,ilocal)+T(itime-  

1,ilocal+1))
enddo
if (my_id.eq.0) then
write(*,*)"Sending and receiving overlap points"
dest = 1
```

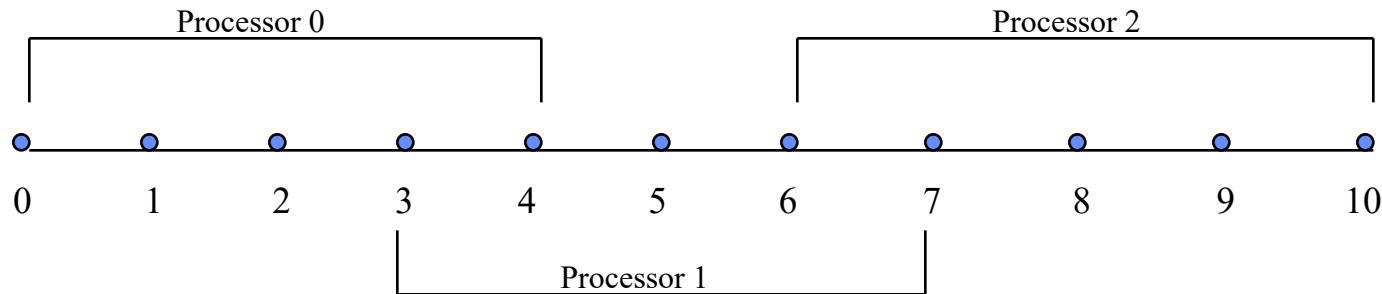
# Fortran MPI Code: 1-D Heat Equation (Contd.)

```
msg_size = 1
    call
MPI_SEND(T(itime,3),msg_size,MPI_DOUBLE_PRECISION,dest,
+         tag,MPI_COMM_WORLD,ierr)
endif
if (my_id.eq.1) then
    from = 0
    dest = 2
    msg_size = 1
    call
MPI_SEND(T(itime,3),msg_size,MPI_DOUBLE_PRECISION,dest,
+         tag,MPI_COMM_WORLD,ierr)
    call
MPI_RECV(T(itime,0),msg_size,MPI_DOUBLE_PRECISION,from
,
+         tag,MPI_COMM_WORLD,status,ierr)
endif
if (my_id.eq.2) then
    from = 1
    dest = 1
    msg_size = 1
    call
MPI_SEND(T(itime,1),msg_size,MPI_DOUBLE_PRECISION,dest,
+         tag,MPI_COMM_WORLD,ierr)
    call
MPI_RECV(T(itime,0),msg_size,MPI_DOUBLE_PRECISION,from
,
+         tag,MPI_COMM_WORLD,status,ierr)
endif
if (my_id.eq.1) then
    from = 2
    dest = 0
    msg_size = 1
    call MPI_RECV(T(itime,4),msg_size,MPI_DOUBLE_PRECISION,from,
+                 tag,MPI_COMM_WORLD,status,ierr)
    call MPI_SEND(T(itime,1),msg_size,MPI_DOUBLE_PRECISION,dest,
+                 tag,MPI_COMM_WORLD,ierr)
endif
if (my_id.eq.0) then
    from = 1
    msg_size = 1
    call MPI_RECV(T(itime,4),msg_size,MPI_DOUBLE_PRECISION,from,
+                 tag,MPI_COMM_WORLD,status,ierr)
endif
enddo

if (my_id.eq.0) then
    write(*,*)"SOLUTION SENT TO FILE AFTER 3 Timesteps:"
endif
FILEN = 'data'//char(my_id+48)//'.dat'
open (5,file=FILEN)
write(5,*)"Processor ",my_id
do ilocal = 0 , 4
    iglobal = 3*my_id + ilocal
    write(5,*)"ilocal=",ilocal,";iglobal=",iglobal,";T=",T(3,ilocal)
enddo
close(5)
call MPI_FINALIZE(ierr)

END
```

# Simple Application using MPI: 1-D Heat Equation



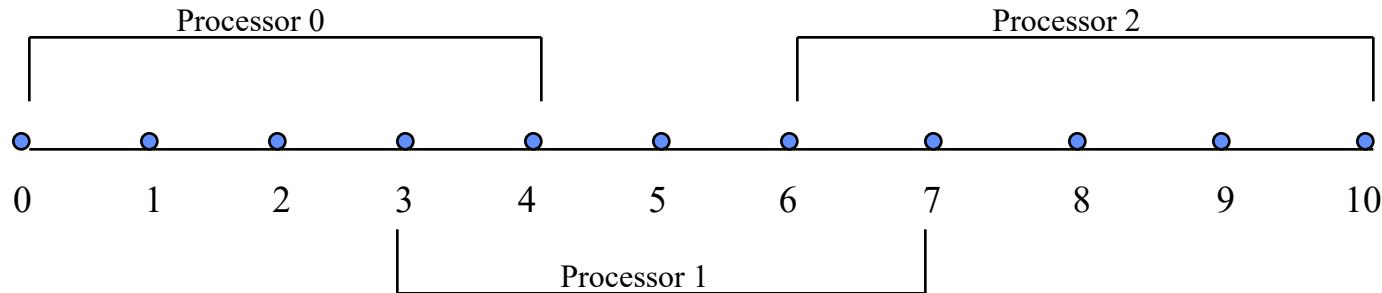
- Compilation

Fortran: `mpif90 -nofree -o heat_mpi.exe heat_mpi.f90`

- Run Job:

`sbatch --res=si2020resday4 heat_mpi.sb`

# Simple Application using MPI: 1-D Heat Equation



## OUTPUT FROM SAMPLE PROGRAM

Process 0 of 3 has started

setting initial conditions

Processor 0 has finished

Process 1 of 3 has started

setting initial conditions

Processor 1 has finished

setting initial conditions

Process 2 of 3 has started

setting initial conditions

Processor 2 has finished

Running Iteration Number 1

setting initial conditions

Sending and receiving overlap points

Running Iteration Number 2

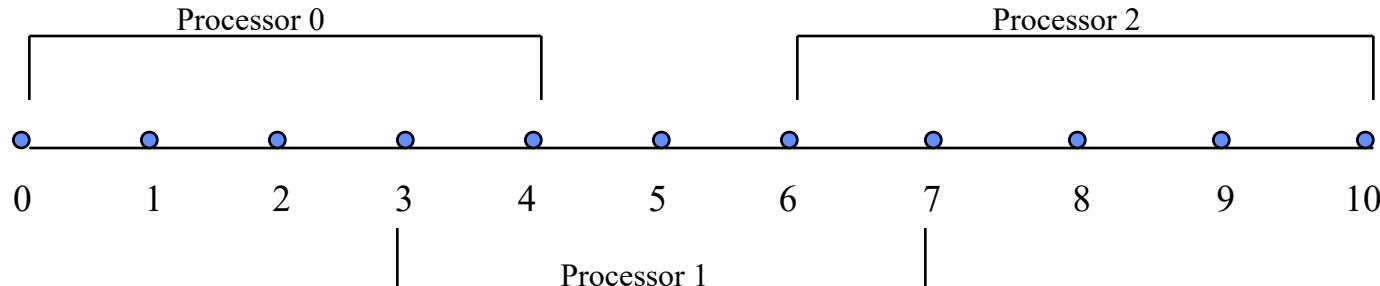
Sending and receiving overlap points

Running Iteration Number 3

Sending and receiving overlap points

SOLUTION SENT TO FILE AFTER 3 Timesteps:

# Simple Application using MPI: 1-D Heat Equation



```
% more data0.dat
```

Processor 0

```
ilocal= 0 ;iglobal= 0 ;T= 0.0000000000000000E+00
ilocal= 1 ;iglobal= 1 ;T= 0.307205621017284991
ilocal= 2 ;iglobal= 2 ;T= 0.584339815421976549
ilocal= 3 ;iglobal= 3 ;T= 0.804274757358271253
ilocal= 4 ;iglobal= 4 ;T= 0.945481682332597884
```

```
% more data2.dat
```

Processor 2

```
ilocal= 0 ;iglobal= 6 ;T= 0.945481682332597995
ilocal= 1 ;iglobal= 7 ;T= 0.804274757358271253
ilocal= 2 ;iglobal= 8 ;T= 0.584339815421976660
ilocal= 3 ;iglobal= 9 ;T= 0.307205621017285102
ilocal= 4 ;iglobal= 10 ;T= 0.0000000000000000E+00
```

```
% more data1.dat
```

Processor 1

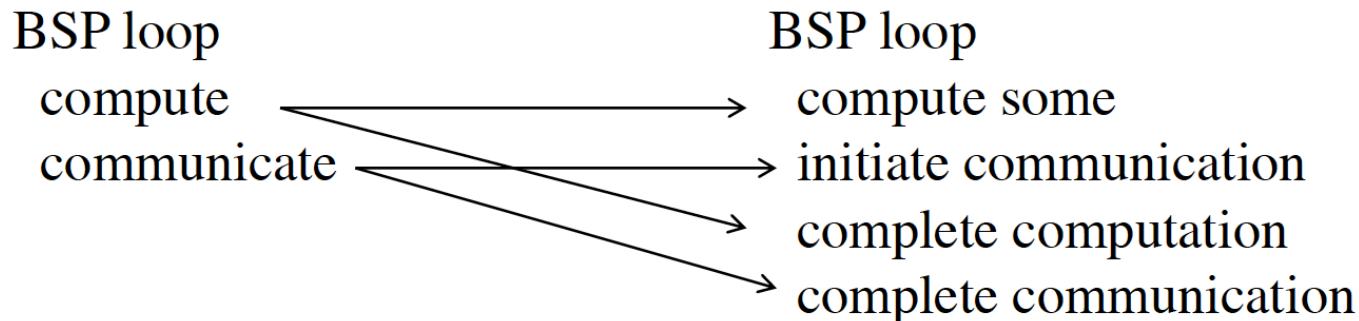
```
ilocal= 0 ;iglobal= 3 ;T= 0.804274757358271253
ilocal= 1 ;iglobal= 4 ;T= 0.945481682332597884
ilocal= 2 ;iglobal= 5 ;T= 0.994138272681972301
ilocal= 3 ;iglobal= 6 ;T= 0.945481682332597995
ilocal= 4 ;iglobal= 7 ;T= 0.804274757358271253
```

# Performance Considerations

- **Overlap communication with computation**
  - Use non-blocking primitives
  - Hide communication cost
  - Split-phase programming
- **Minimize surface-to-volume ratio**
  - Ghost cell exchange
- **Avoid communication**
  - Even at the cost of some more computation
  - Example: double size of ghost cell and communicate every other time step
  - Communication avoiding algorithms

# Asynchronous Communication

- **Overlap communication w/ computation**
  - High performance interconnects can offload communication tasks from CPU to adapter
- **Condition**
  - No data dependencies on transfer
- **Split-phase programming**



# MPI – Profiling, Tracing Tools

- Several options available. On Comet we have mpiP, TAU, and Allinea MAP installed.
- Useful when you are trying to isolate performance issues.
- Tools can give you info on how much time is being spent in communication. The levels of detail vary with each tool.
- In general identify scaling bottlenecks and try to overlap communication with computation where possible.

# mpiP example

- Location: \$HOME/PARALLEL/MISC

- Compile:

```
mpif90 -nofree -g -o heat_mpi_profile.exe heat_mpi.f90 -  
L/share/apps/compute/mpiP/v3.4.1/mv2/lib -lmpiP -L/share/apps/compute/libiberty -liberty -  
L/share/apps/compute/libunwind/v1.1/mv2/lib -lunwind
```

- Executable already exists. Just submit  
**heat\_mpi\_profile.sb**.
- Once the job runs you get a .mpiP file.

# mpiP output

```
@ Command : ./heat_mpi_profile.exe
@ Version          : 3.4.1
@ MPIP Build date : Aug  9 2018, 12:07:41
@ Start time       : 2018 08 09 12:12:23
@ Stop time        : 2018 08 09 12:12:23
@ Timer Used      : PMPI_Wtime
@ MPIP env var    : [null]
@ Collector Rank   : 0
@ Collector PID    : 31588
@ Final Output Dir: .
@ Report generation: Single collector task
@ MPI Task Assignment: 0 comet-06-13.sdsc.edu
@ MPI Task Assignment: 1 comet-06-13.sdsc.edu
@ MPI Task Assignment: 2 comet-06-13.sdsc.edu
```

---

```
-----  
@--- MPI Time (seconds) -----  
-----
```

Task	AppTime	MPITime	MPI%
0	0.024	0.00989	41.22
1	0.0243	0.0099	40.81
2	0.0243	0.000126	0.52

# mpiP Output

```
*      0.0726    0.0199   27.44
-----
@--- Callsites: 1 -----
ID Lev File/Address          Line Parent_Funct           MPI_Call
 1  0 0x40ca45              [unknown]                 Recv
-----
@--- Aggregate Time (top twenty, descending, milliseconds) -----
Call             Site     Time    App%    MPI%    COV
Recv            1       19.5   26.87   97.93   0.86
Send            1       0.412   0.57    2.07    0.21
-----
@--- Aggregate Sent Message Size (top twenty, descending, bytes) -----
Call             Site     Count   Total    Avrg    Sent%
Send            1       12      96        8       100.00
-----
@--- Callsite Time statistics (all, milliseconds): 6 -----
Name            Site Rank  Count   Max     Mean    Min     App%    MPI%
Recv            1    0     3      9.73   3.25   0.00504  40.63  98.55
Recv            1    1     6      9.69   1.62   0.00194  40.14  98.36
```

# mpiP output

```
-----  
@---- Callsite Time statistics (all, milliseconds): 6 -----  
  
Name      Site Rank Count   Max     Mean    Min   App%  MPI%  
Recv       1     0     3     9.73    3.25  0.00504  40.63  98.55  
Recv       1     1     6     9.69    1.62  0.00194  40.14  98.36  
Recv       1     2     3     0.015   0.00664  0.00188  0.08  15.83  
Send       1     0     3     0.136   0.0477  0.00303  0.60  1.45  
Send       1     1     6     0.146   0.0272  0.00203  0.67  1.64  
Send       1     2     3     0.0999  0.0353  0.00298  0.44  84.17  
Send       1     *    24     9.73    0.83  0.00188  27.44 100.00  
-----  
@---- Callsite Message Sent statistics (all, sent bytes) -----  
  
Name      Site Rank Count   Max     Mean    Min   Sum  
Send       1     0     3      8      8      8    24  
Send       1     1     6      8      8      8    48  
Send       1     2     3      8      8      8    24  
Send       1     *    12     8      8      8    96  
-----  
@---- End of Report -----  
-----
```

# More Complex routines

- Derived Data Types
- User defined reduction functions
- Groups/communicator management
- Parallel I/O
- One Sided Communication Routines (RDMA)
- MPI-3 Standard has over 400 routines(!).

# Homework!

- **Change directory to \$HOME/PARALLEL/MISC**
- **Code “sample.f” has two bugs.**
- **Compile:**

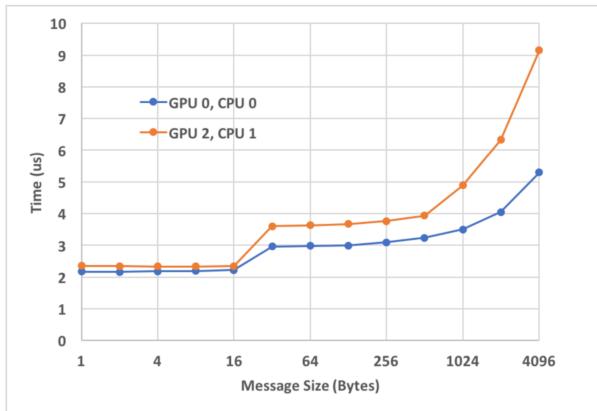
```
mpif90 -o sample.exe sample.f
```

Run: sbatch sample.sb

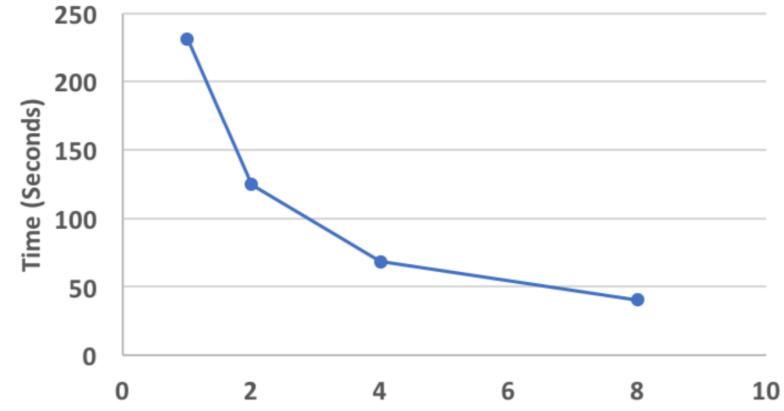
See if you can identify the two bugs!

# MVAPICH2-GDR on Comet

- Dr. DK Panda's team (Ohio State), have been enhancing MVAPICH2 to utilize GPUDirect capabilities and achieve low MPI latencies and high bandwidth.
- MVAPICH2-GDR is available on the SDSC Comet GPU nodes featuring NVIDIA K80s and P100s.
- Applications that use MVAPICH2-GDR include HOOMD-Blue, OSU-Caffe, and TensorFlow (Horovod).



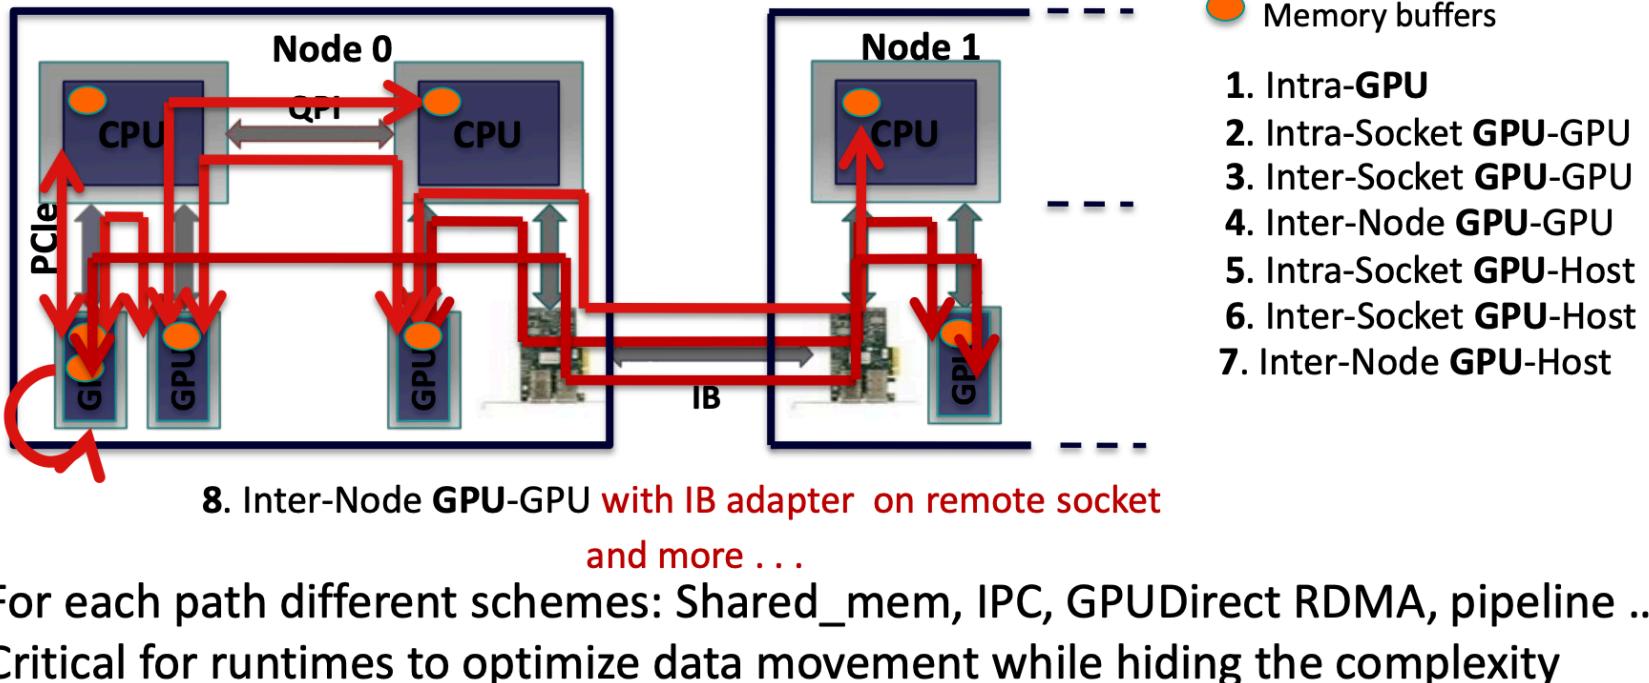
- Latency between GPU 0 , process bound to CPU 0 on both nodes: 2.17  $\mu$ s
- Latency between GPU 2 , process bound to CPU 1 on both nodes: 2.35  $\mu$ s



Scaling of OSU-Caffe performance w.r.t number of  
GPUS; Runs using MVAPICH2-GDR

# Optimizing MPI Data Movement on GPU Clusters

- Connected as PCIe devices – Flexibility but Complexity



- For each path different schemes: Shared\_mem, IPC, GPUDirect RDMA, pipeline ...
- Critical for runtimes to optimize data movement while hiding the complexity

Reference: Dhabaleswar K. (DK) Panda, Hari Subramoni “*MVAPICH2-GDR: Pushing the Frontier of HPC and Deep Learning*”, GTC 2017

# GPU-Aware (CUDA-Aware) MPI Library: MVAPICH2-GPU

- Standard MPI interfaces used for unified data movement
- Takes advantage of Unified Virtual Addressing (>= CUDA 4.0)
- Overlaps data movement from GPU with RDMA transfers

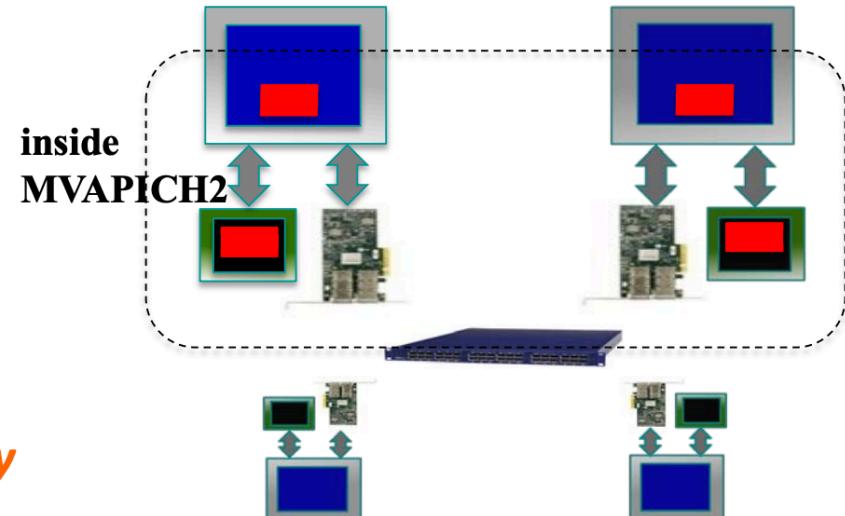
## At Sender:

```
MPI_Send(s_devbuf, size, ...);
```

## At Receiver:

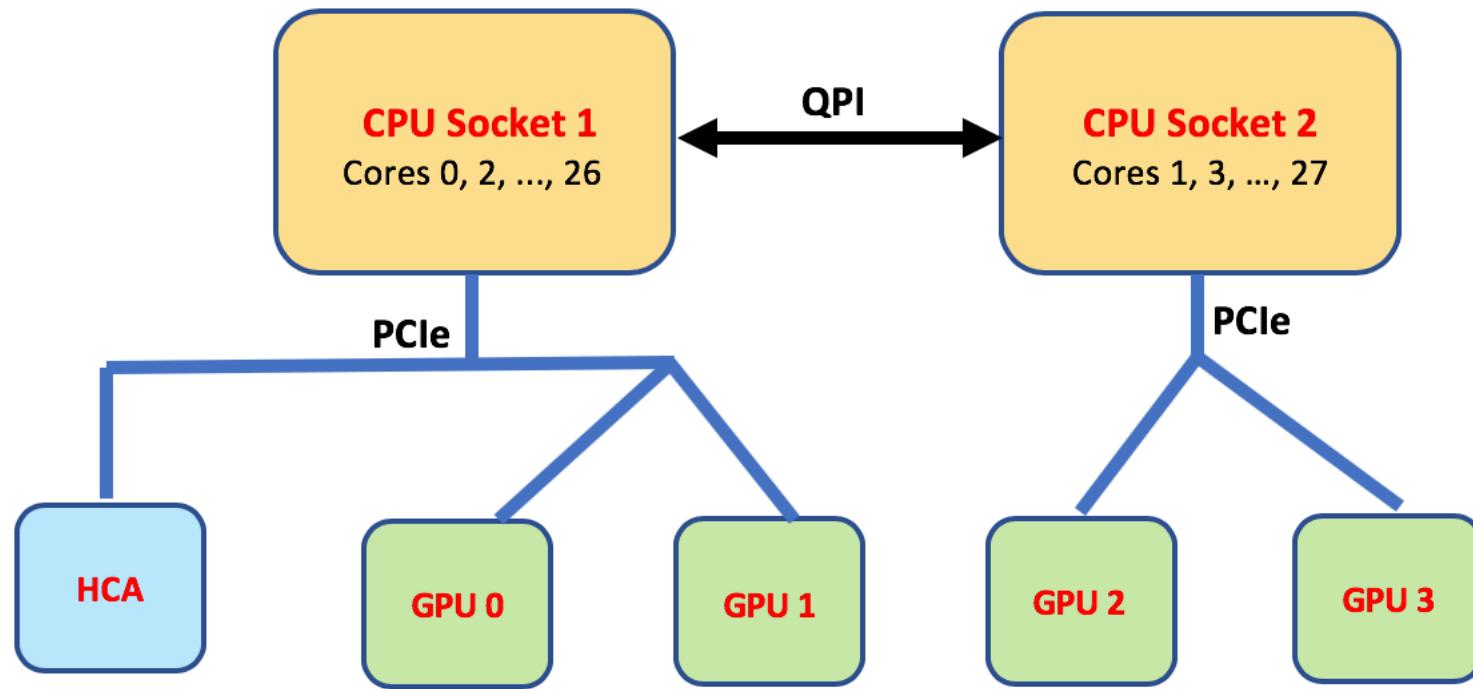
```
MPI_Recv(r_devbuf, size, ...);
```

***High Performance and High Productivity***



# MVAPICH2-GDR

## Comet P100 node layout



# Comet K80 node architecture

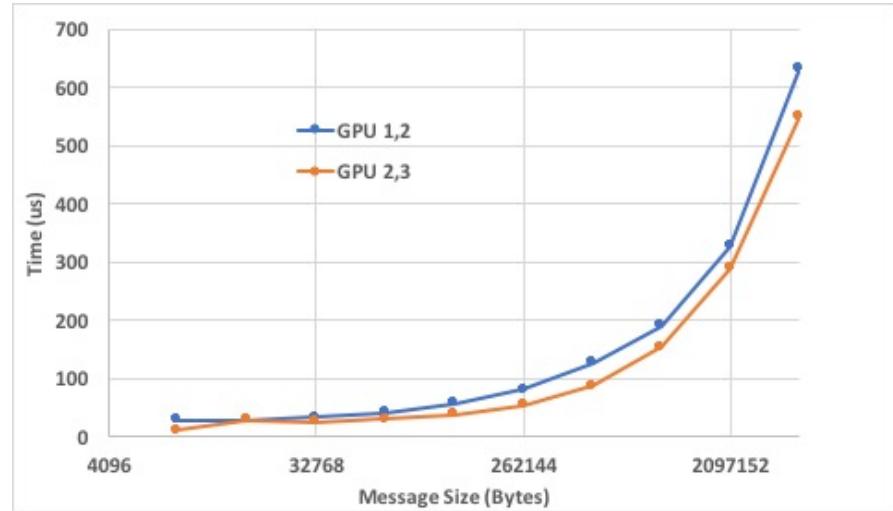
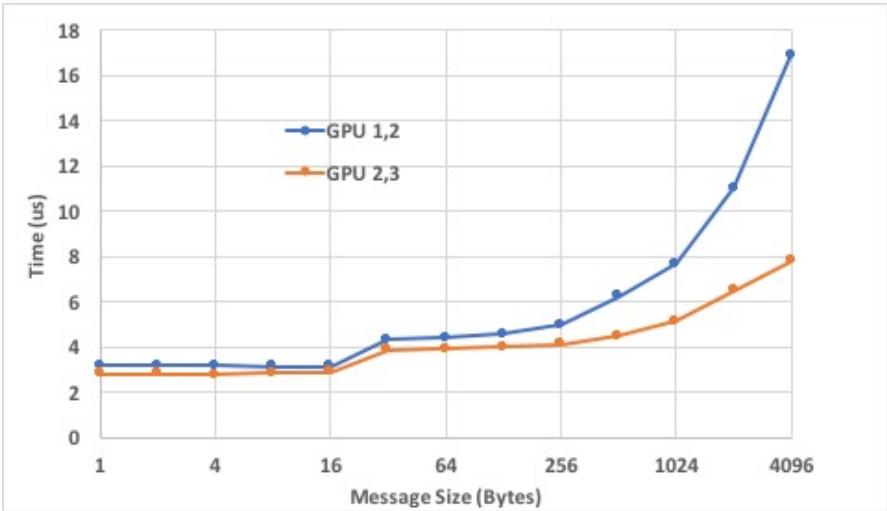
	GPU0	GPU1	GPU2	GPU3	mlx4_0	CPU Affinity
GPU0	X	PIX	SOC	SOC	SOC	0-0,2-2,4-4,6-6,8-8,10-10,12-12,14-14,16-16,18-18,20-20,22-22
GPU1	PIX	X	SOC	SOC	SOC	0-0,2-2,4-4,6-6,8-8,10-10,12-12,14-14,16-16,18-18,20-20,22-22
GPU2	SOC	SOC	X	PIX	PHB	1-1,3-3,5-5,7-7,9-9,11-11,13-13,15-15,17-17,19-19,21-21,23-23
GPU3	SOC	SOC	PIX	X	PHB	1-1,3-3,5-5,7-7,9-9,11-11,13-13,15-15,17-17,19-19,21-21,23-23
mlx4_0	SOC	SOC	PHB	PHB	X	

Legend:

X = Self  
SOC = Connection traversing PCIe as well as the SMP link between CPU sockets(e.g. QPI)  
PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)  
PXB = Connection traversing multiple PCIe switches (without traversing the PCIe Host Bridge)  
PIX = Connection traversing a single PCIe switch  
NV# = Connection traversing a bonded set of # NVLinks

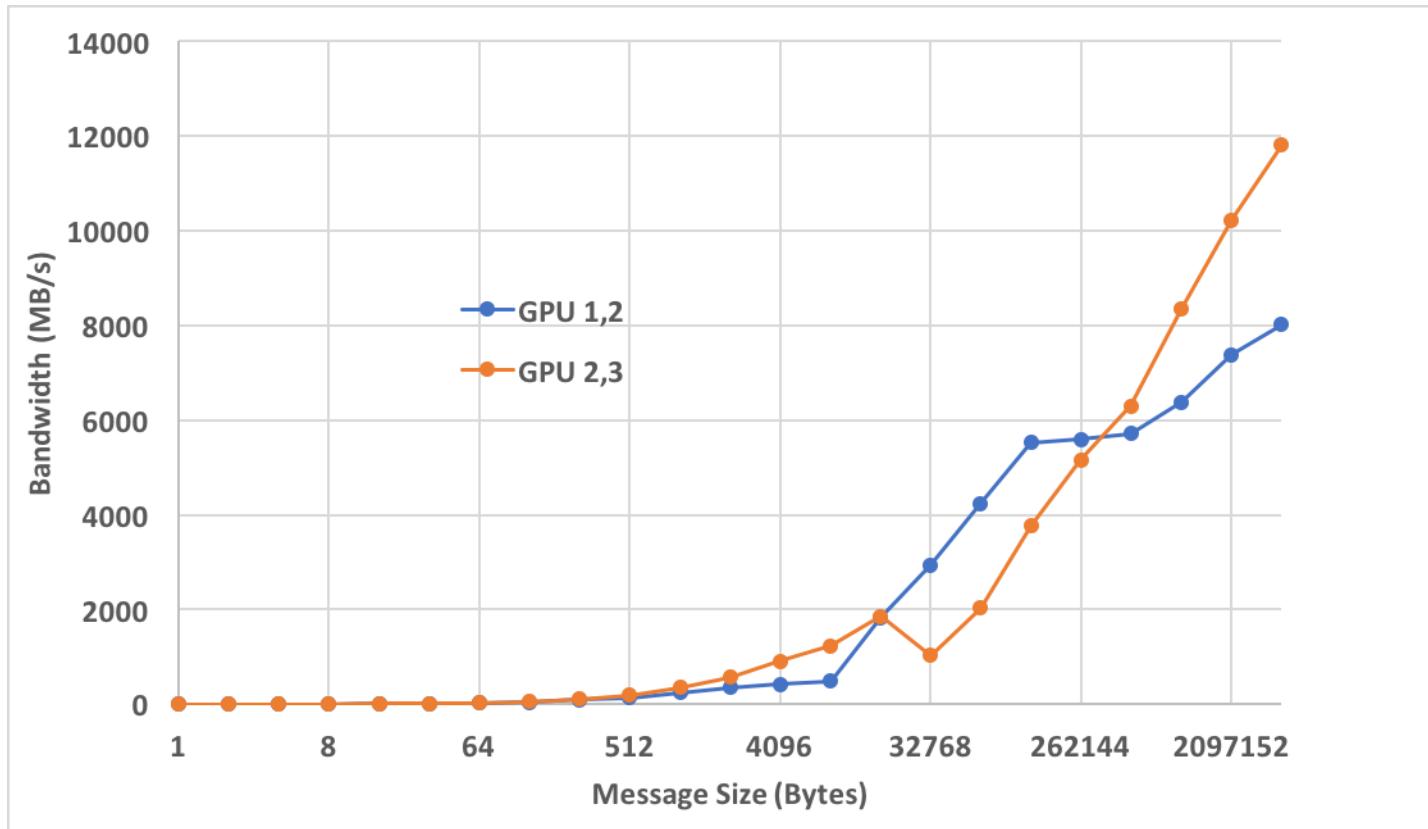
- 4 GPUs per node
- GPUs (0,1) and (2,3) can do P2P communication
- Mellanox InfiniBand adapter associated with second socket (GPUs 2, 3)

# OSU Latency (osu\_latency) Benchmark Intra-node, K80 nodes

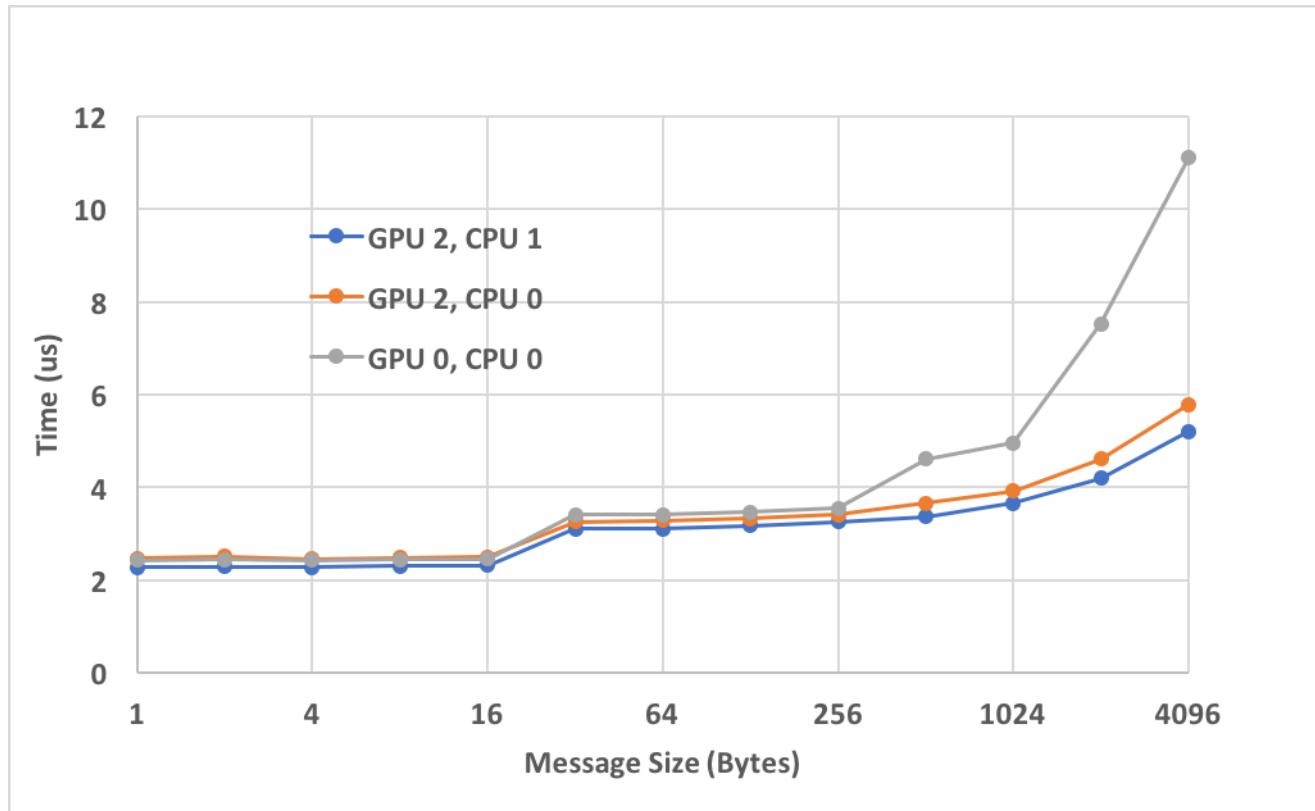


- Latency between GPU 2 , GPU 3:  $2.82 \mu\text{s}$
- Latency between GPU 1 , GPU 2:  $3.18 \mu\text{s}$

# OSU Bandwidth (osu\_bw) Benchmark Intra-node, K80 nodes

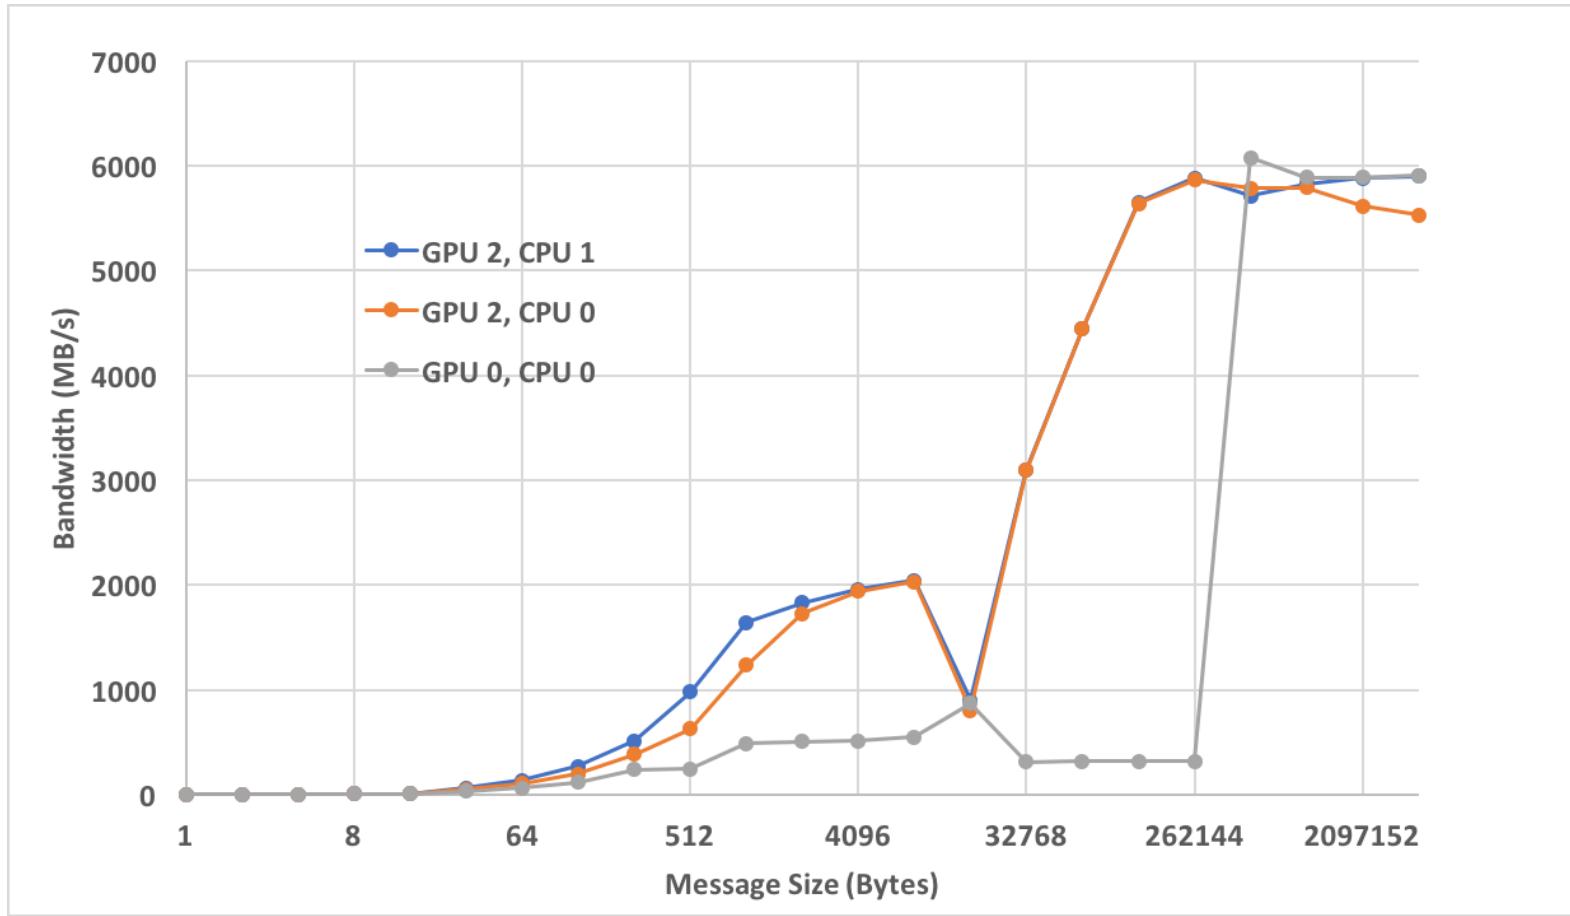


# OSU Latency (osu\_latency) Benchmark Inter-node, K80 nodes



- Latency between GPU 2 , process bound to CPU 1 on both nodes:  $2.27 \mu\text{s}$
- Latency between GPU 2 , process bound to CPU 0 on both nodes:  $2.47 \mu\text{s}$
- Latency between GPU 0 , process bound to CPU 0 on both nodes:  $2.43 \mu\text{s}$

# OSU Bandwidth (osu\_bw) Benchmark Inter-node, K80 nodes



# Comet P100 node architecture

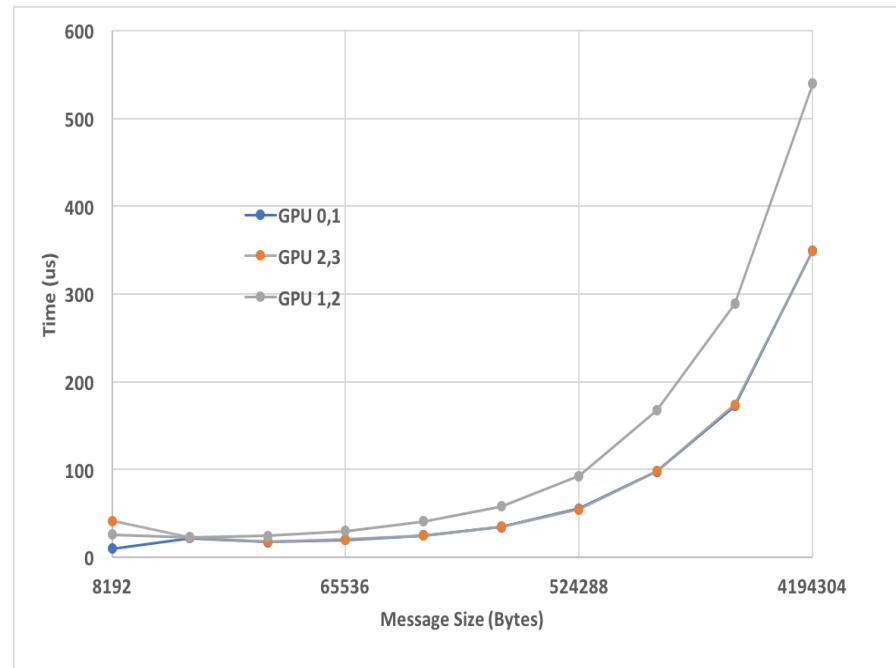
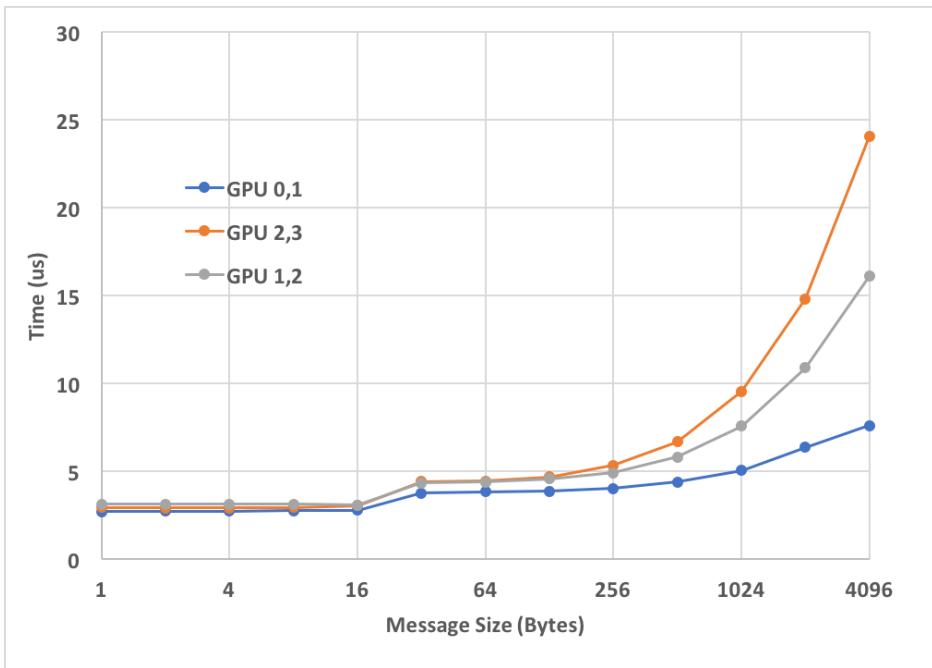
	GPU0	GPU1	GPU2	GPU3	mlx4_0	CPU Affinity
GPU0	X	PIX	SOC	SOC	PHB	0-0,2-2,4-4,6-6,8-8,10-10,12-12,14-14,16-16,18-18,20-20,22-22,24-24,26-26
GPU1	PIX	X	SOC	SOC	PHB	0-0,2-2,4-4,6-6,8-8,10-10,12-12,14-14,16-16,18-18,20-20,22-22,24-24,26-26
GPU2	SOC	SOC	X	PIX	SOC	1-1,3-3,5-5,7-7,9-9,11-11,13-13,15-15,17-17,19-19,21-21,23-23,25-25,27-27
GPU3	SOC	SOC	PIX	X	SOC	1-1,3-3,5-5,7-7,9-9,11-11,13-13,15-15,17-17,19-19,21-21,23-23,25-25,27-27
mlx4_0	PHB	PHB	SOC	SOC	X	

Legend:

X = Self  
SOC = Connection traversing PCIe as well as the SMP link between CPU sockets(e.g. QPI)  
PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)  
PXB = Connection traversing multiple PCIe switches (without traversing the PCIe Host Bridge)  
PIX = Connection traversing a single PCIe switch  
NV# = Connection traversing a bonded set of # NVLinks

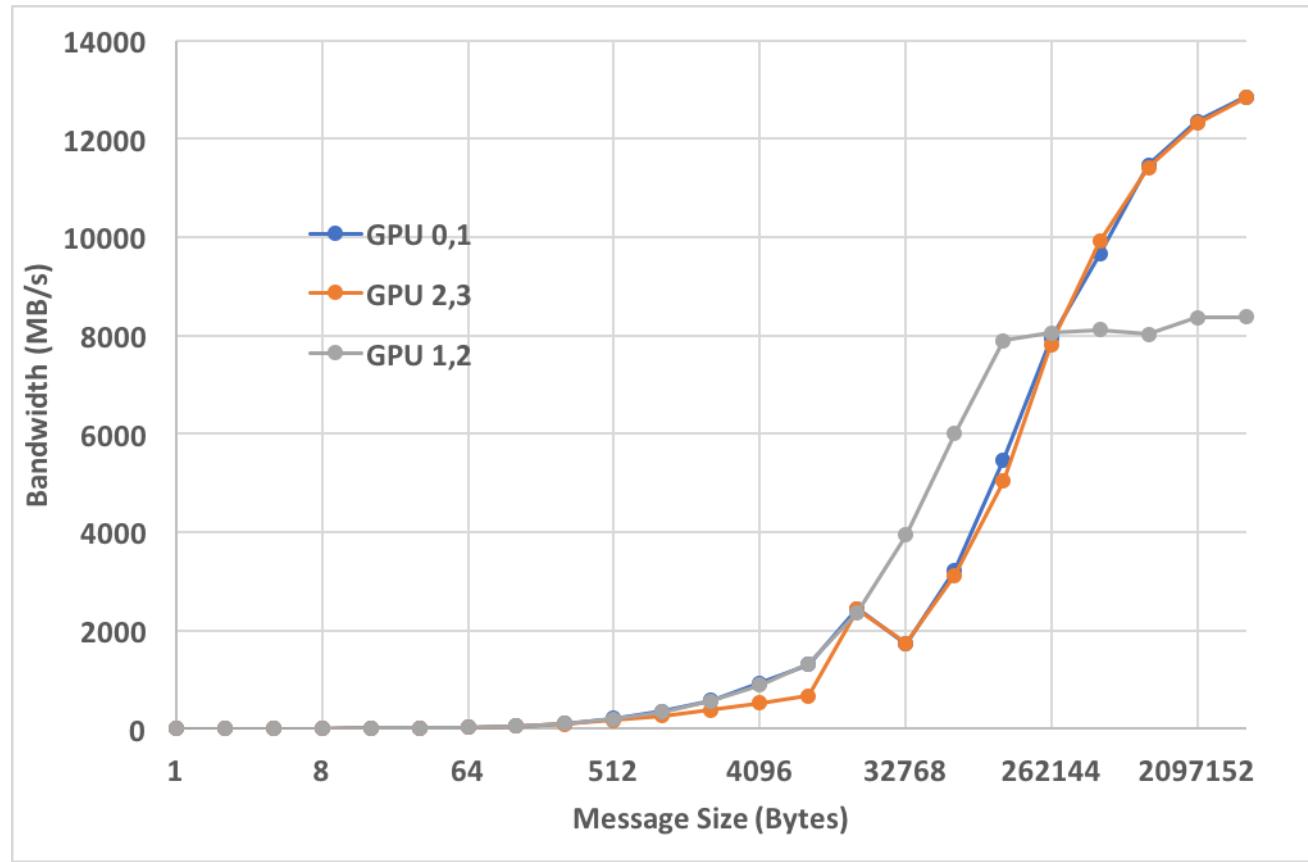
- 4 GPUs per node
- GPUs (0,1) and (2,3) can do P2P communication
- Mellanox InfiniBand adapter associated with first socket (GPUs 0, 1)

# OSU Latency (osu\_latency) Benchmark Intra-node, P100 nodes

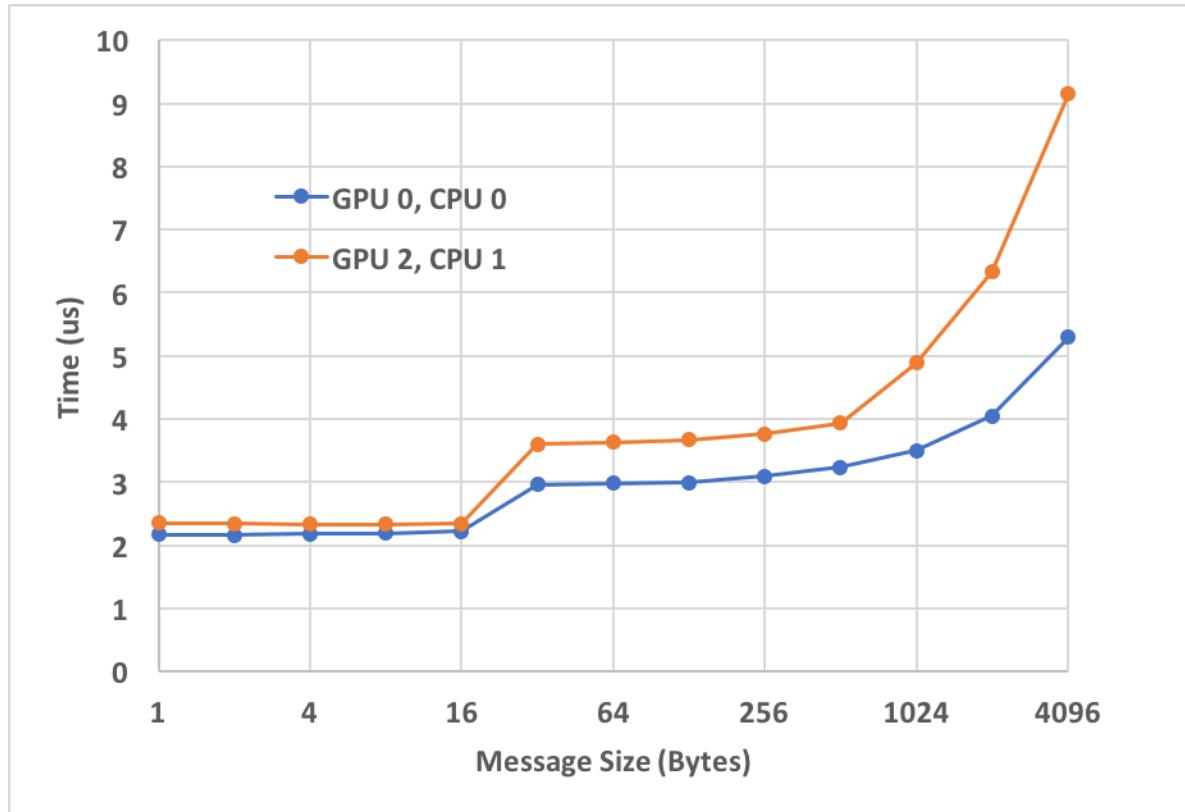


- Latency between GPU 0 , GPU 1:  $2.73 \mu\text{s}$
- Latency between GPU 2 , GPU 3:  $2.95 \mu\text{s}$
- Latency between GPU 1 , GPU 2:  $3.13 \mu\text{s}$

# OSU Bandwidth (osu\_bw) Benchmark Intra-node, P100 nodes



# OSU Latency (`osu_latency`) Benchmark Inter-node, P100 nodes



- Latency between GPU 0 , process bound to CPU 0 on both nodes:  $2.17 \mu\text{s}$
- Latency between GPU 2 , process bound to CPU 1 on both nodes:  $2.35 \mu\text{s}$

# References

- **Excellent tutorials from LLNL:**
  - <https://computing.llnl.gov/tutorials/mpi/>
  - <https://computing.llnl.gov/tutorials/openMP/>
- **MPI for Python:**
  - <http://mpi4py.scipy.org/docs/usrman/tutorial.html>
- **MVAPICH2 User Guide:**
  - <http://mvapich.cse.ohio-state.edu/userguide/>
- **MVAPICH2-GDR User Guide:**
  - <http://mvapich.cse.ohio-state.edu/userguide/gdr/>